

An aerial photograph of the Colorado School of Mines campus. In the foreground, there are several large, modern buildings with light-colored facades and flat roofs. A large green field, possibly a sports field, is visible in the middle ground. In the background, there are rolling hills with sparse vegetation and some residential houses. The sky is blue with a few wispy clouds.

Colorado School of Mines

Computer Vision

Professor William Hoff

Dept of Electrical Engineering & Computer Science


<http://inside.mines.edu/~whoff/>

3D-3D Coordinate Transforms

An excellent reference is the book “Introduction to Robotics” by John Craig

3D Coordinate Systems



- Coordinate frames
 - Denote as $\{A\}$, $\{B\}$, etc
 - Examples: camera, world, model
- The pose* of $\{B\}$ with respect to $\{A\}$ is described by
 - Translation, or position
 - Rotation, or orientation
- Translation is just a 3D vector; rotation is more complicated




**The term “pose” means position and orientation*

Rotations in 3D

- A 3D rotation has 3 degrees of freedom
 - Namely, it takes 3 numbers to describe the orientation of an object in the world
 - Think of “roll”, “pitch”, “yaw” for an airplane
- We can represent a 3D rotation by doing successive rotations about the X,Y, and Z axes



You are probably familiar with rotations in 2D



It can be represented by a 2x2 rotation matrix


$$\mathbf{R} = \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix}$$

$$\begin{pmatrix} x \\ y \end{pmatrix} = \mathbf{R} \begin{pmatrix} x' \\ y' \end{pmatrix}$$

- Note: \mathbf{R} is orthonormal
 - Rows, columns are orthogonal ($\mathbf{r}_1 \cdot \mathbf{r}_2 = 0$, $\mathbf{c}_1 \cdot \mathbf{c}_2 = 0$)
 - Transpose is the inverse; $\mathbf{R}\mathbf{R}^T = \mathbf{I}$
 - Determinant is $|\mathbf{R}| = 1$

To do 3D rotations, we can create 2D rotation matrices for each axis


- Rotation about the Z axis



⊙ Points toward me
⊗ Points away from me


$$\begin{pmatrix} {}^B x \\ {}^B y \\ {}^B z \end{pmatrix} = \begin{pmatrix} \cos \theta_Z & -\sin \theta_Z & 0 \\ \sin \theta_Z & \cos \theta_Z & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} {}^A x \\ {}^A y \\ {}^A z \end{pmatrix}$$

- Rotation about the X axis



$$\begin{pmatrix} {}^B x \\ {}^B y \\ {}^B z \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \cos \theta_X & -\sin \theta_X \\ 0 & \sin \theta_X & \cos \theta_X \end{pmatrix} \begin{pmatrix} {}^A x \\ {}^A y \\ {}^A z \end{pmatrix}$$

- Rotation about the Y axis



$$\begin{pmatrix} {}^B x \\ {}^B y \\ {}^B z \end{pmatrix} = \begin{pmatrix} \cos \theta_Y & 0 & \sin \theta_Y \\ 0 & 1 & 0 \\ -\sin \theta_Y & 0 & \cos \theta_Y \end{pmatrix} \begin{pmatrix} {}^A x \\ {}^A y \\ {}^A z \end{pmatrix}$$

Note signs are different than in the other cases

3D Rotation Matrix

- We can concatenate the 3 rotations to yield a single 3x3 rotation matrix; e.g.,

$$\begin{aligned} {}^A_B R_{XYZ}(\theta_X, \theta_Y, \theta_Z) &= R_Z(\theta_Z) R_Y(\theta_Y) R_X(\theta_X) \\ &= \begin{pmatrix} c_Z & -s_Z & 0 \\ s_Z & c_Z & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} c_Y & 0 & s_Y \\ 0 & 1 & 0 \\ -s_Y & 0 & c_Y \end{pmatrix} \begin{pmatrix} 1 & 0 & 0 \\ 0 & c_X & -s_X \\ 0 & s_X & c_X \end{pmatrix} \end{aligned}$$

where

$$c_X = \cos(\theta_X), s_Y = \sin(\theta_Y), \text{ etc}$$

- Note: we use the convention that to rotate a vector, we pre-multiply it; i.e., $\mathbf{v}' = \mathbf{R} \mathbf{v}$
 - This means that if $\mathbf{R} = \mathbf{R}_Z \mathbf{R}_Y \mathbf{R}_X$, we actually apply the X rotation first, then the Y rotation, then the Z rotation

Python: Creating a Rotation Matrix

```
import numpy as np
```

```
ax, ay, az = 0.1, -0.2, 0.3 # radians
```

```
sx, sy, sz = np.sin(ax), np.sin(ay), np.sin(az)  
cx, cy, cz = np.cos(ax), np.cos(ay), np.cos(az)
```

```
Rx = np.array(((1, 0, 0), (0, cx, -sx), (0, sx, cx)))  
Ry = np.array(((cy, 0, sy), (0, 1, 0), (-sy, 0, cy)))  
Rz = np.array(((cz, -sz, 0), (sz, cz, 0), (0, 0, 1)))
```

*Note the strange symbol @
for matrix multiplication!*

```
# Apply X rotation first, then Y, then Z  
R = Rz @ Ry @ Rx # Use @ for matrix mult  
print(R)
```

```
# Apply Z rotation first, then Y, then X  
R = Rx @ Ry @ Rz  
print(R)
```


Python: Creating a Rotation Matrix

```
import numpy as np
```

```
ax, ay, az = 0.1, -0.2, 0.3 # radians
```

```
sx, sy, sz = np.sin(ax), np.sin(ay), np.sin(az)  
cx, cy, cz = np.cos(ax), np.cos(ay), np.cos(az)
```

```
Rx = np.array(((1, 0, 0), (0, cx, -sx), (0, sx, cx)))  
Ry = np.array(((cy, 0, sy), (0, 1, 0), (-sy, 0, cy)))  
Rz = np.array(((cz, -sz, 0), (sz, cz, 0), (0, 0, 1)))
```

```
# Apply X rotation first, then Y, then Z
```

```
R = Rz @ Ry @ Rx # Use @ for matrix mult
```

```
print(R)
```

```
[[ 0.93629336 -0.31299183 -0.15934508]  
 [ 0.28962948  0.94470249 -0.153792  ]  
 [ 0.19866933  0.0978434  0.97517033]]
```

```
# Apply Z rotation first, then Y, then X
```

```
R = Rx @ Ry @ Rz
```

```
print(R)
```

Python: Creating a Rotation Matrix

```
import numpy as np
```

```
ax, ay, az = 0.1, -0.2, 0.3 # radians
```

```
sx, sy, sz = np.sin(ax), np.sin(ay), np.sin(az)
```

```
cx, cy, cz = np.cos(ax), np.cos(ay), np.cos(az)
```

```
Rx = np.array(((1, 0, 0), (0, cx, -sx), (0, sx, cx)))
```

```
Ry = np.array(((cy, 0, sy), (0, 1, 0), (-sy, 0, cy)))
```

```
Rz = np.array(((cz, -sz, 0), (sz, cz, 0), (0, 0, 1)))
```

```
# Apply X rotation first, then Y, then Z
```

```
R = Rz @ Ry @ Rx # Use @ for matrix mult
```

```
print(R)
```

```
[[ 0.93629336 -0.31299183 -0.15934508]
 [ 0.28962948  0.94470249 -0.153792  ]
 [ 0.19866933  0.0978434  0.97517033]]
```

```
# Apply Z rotation first, then Y, then X
```

```
R = Rx @ Ry @ Rz
```

```
print(R)
```

```
[[ 0.93629336 -0.28962948 -0.19866933]
 [ 0.27509585  0.95642509 -0.0978434 ]
 [ 0.21835066  0.03695701  0.97517033]]
```

Different!



Problems with XYZ angles

- XYZ angles are intuitive, but they are not good for computation 😞
 - The result depends on the order in which the transforms are applied; i.e., XYZ or ZYX
 - Sometimes one or more angles change dramatically in response to a small change in orientation
 - Some orientations have singularities; i.e., the angles are not well defined
- We'll just use XYZ angles to create a rotation matrix, then work with the rotation matrix
- The rotation matrix is always unique for a given orientation


3D Rotation Matrix as a Transformation

- \mathbf{R} can represent a rotational transformation of one frame to another
- We can rotate a vector represented in frame A to obtain its representation in frame B

$${}^B \mathbf{v} = {}^B \mathbf{R} {}^A \mathbf{v}$$

- Note: as in 2D, rotation matrices are orthonormal so the inverse of a rotation matrix is just its transpose

$${}^B \mathbf{R} = \begin{pmatrix} r_{11} & r_{12} & r_{13} \\ r_{21} & r_{22} & r_{23} \\ r_{31} & r_{32} & r_{33} \end{pmatrix}$$



$$\left({}^B \mathbf{R}\right)^{-1} = \left({}^B \mathbf{R}\right)^T = {}^A \mathbf{R}$$

Notation

- For vectors, such as ${}^A\mathbf{v}$, the leading superscript represents the coordinate frame that the vector is expressed in
- For transforms, such as ${}^B_A\mathbf{R}$, this matrix represents a rotational transformation of frame A to frame B
 - The leading subscript indicates “from”
 - The leading superscript indicates “to”

$${}^A\mathbf{v} = \begin{pmatrix} {}^Ax \\ {}^Ay \\ {}^Az \end{pmatrix}$$

3D Rotation Matrix

- The elements of \mathbf{R} are direction cosines (the projections of unit vectors from one frame onto the unit vectors of the other frame)
- The columns of \mathbf{R} are the unit vectors of A, expressed in the B frame
- The rows of \mathbf{R} are the unit vectors of {B} expressed in {A}

$${}^B_A \mathbf{R} = \begin{pmatrix} \hat{\mathbf{x}}_A \cdot \hat{\mathbf{x}}_B & \hat{\mathbf{y}}_A \cdot \hat{\mathbf{x}}_B & \hat{\mathbf{z}}_A \cdot \hat{\mathbf{x}}_B \\ \hat{\mathbf{x}}_A \cdot \hat{\mathbf{y}}_B & \hat{\mathbf{y}}_A \cdot \hat{\mathbf{y}}_B & \hat{\mathbf{z}}_A \cdot \hat{\mathbf{y}}_B \\ \hat{\mathbf{x}}_A \cdot \hat{\mathbf{z}}_B & \hat{\mathbf{y}}_A \cdot \hat{\mathbf{z}}_B & \hat{\mathbf{z}}_A \cdot \hat{\mathbf{z}}_B \end{pmatrix}$$

$${}^B_A \mathbf{R} = \left(\begin{pmatrix} {}^B \hat{\mathbf{x}}_A \\ {}^B \hat{\mathbf{y}}_A \\ {}^B \hat{\mathbf{z}}_A \end{pmatrix} \right)$$

$${}^B_A \mathbf{R} = \left(\begin{pmatrix} {}^A \hat{\mathbf{x}}_B^T \\ {}^A \hat{\mathbf{y}}_B^T \\ {}^A \hat{\mathbf{z}}_B^T \end{pmatrix} \right)$$

$${}^B_A \mathbf{R} {}^A \hat{\mathbf{x}}_A = \begin{pmatrix} r_{11} & r_{12} & r_{13} \\ r_{21} & r_{22} & r_{23} \\ r_{31} & r_{32} & r_{33} \end{pmatrix} \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} = \begin{pmatrix} r_{11} \\ r_{21} \\ r_{31} \end{pmatrix} = {}^B \hat{\mathbf{x}}_A$$

Doing a rotation AND a translation


- We can use \mathbf{R}, \mathbf{t} to transform a point from coordinate frame $\{B\}$ to frame $\{A\}$

$${}^A\mathbf{p} = {}^A_B\mathbf{R} {}^B\mathbf{p} + \mathbf{t}$$

- Where
 - ${}^A\mathbf{p}$ is the representation of \mathbf{p} in frame $\{A\}$
 - ${}^B\mathbf{p}$ is the representation of \mathbf{p} in frame $\{B\}$

- Note

\mathbf{t} is the translation of B's origin in the A frame, ${}^A\mathbf{t}_{Borg}$



Homogeneous Coordinates

- We can represent the transformation with a *single* matrix multiplication if we write \mathbf{p} in *homogeneous* coordinates
 - This simply means to append a 1 as a 4th element
 - If the 4th element ever becomes $\neq 1$, we divide through by it

The leading superscript indicates what coordinate frame the point is represented in

→ ${}^A\mathbf{p} = \begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix} = \begin{pmatrix} sx \\ sy \\ sz \\ s \end{pmatrix}$

- Then

$${}^B\mathbf{p} = \mathbf{H} {}^A\mathbf{p} \quad \text{where} \quad \mathbf{H} = \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

General Rigid Transformation

- A general rigid transformation can be represented by a single 4x4 homogeneous transformation matrix

$${}^A_B\mathbf{H} = \begin{bmatrix} {}^A_B\mathbf{R} & {}^A\mathbf{t}_{Borg} \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- A note on notation: ${}^A\mathbf{p} = {}^A_B\mathbf{H} {}^B\mathbf{p}$
 Cancel leading subscript with trailing superscript

- Can concatenate transformations together
 - Leading subscripts cancel trailing superscripts

$${}^C_A\mathbf{H} = {}^C_B\mathbf{H} {}^B_A\mathbf{H} \quad {}^D_A\mathbf{H} = {}^D_C\mathbf{H} {}^C_B\mathbf{H} {}^B_A\mathbf{H}, \quad \text{etc}$$

Example: Transforming a point, using Python

- Assume we have a point $\mathbf{p} = (-1,0,1)^T$, in frame A
- Transform it to frame B, if the pose of frame B with respect to frame A is given by ${}^B_A\mathbf{R}$ and ${}^B\mathbf{t}_{Aorg}$

=> We need to do ${}^B\mathbf{p} = {}^B_A\mathbf{H} {}^A\mathbf{p}$

Rotation matrix of A with respect to B.

`R_A_B = np.array(((1,0,0),(0,0,-1),(0,1,0)))` *# Get 3x3 matrix*

The translation is the origin of A in B.

`tAorg_B = np.array([[1,2,4]]).T` *# Get as a 3x1 matrix*

H_A_B means transform A to B.

`H_A_B = np.block([R_A_B, tAorg_B], [0,0,0,1])` *# Get 4x4 matrix*

Define a point in the A frame, as [x,y,z,1].

`P_A = np.array([[-1,0,1,1]]).T` *# Get as a 4x1 matrix*

Convert point to B frame.

`P_B = H_A_B @ P_A`

Assume

$${}^B_A\mathbf{R} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 0 & -1 \\ 0 & 1 & 0 \end{pmatrix}$$

$${}^B\mathbf{t}_{Aorg} = \begin{pmatrix} 1 \\ 2 \\ 4 \end{pmatrix}$$

Inverse Transformations

- The **matrix inverse** is the inverse transformation



$${}^A_B\mathbf{H} = \left({}^B_A\mathbf{H}\right)^{-1}$$

- Note – unlike rotation matrices, the inverse of a full 4x4 homogeneous transformation matrix is **not the transpose**

$${}^A_B\mathbf{H} \neq \left({}^B_A\mathbf{H}\right)^T$$

Example

- A camera is located at point $(0, -5, 3)$ with respect to the world. The camera is tilted down by 30 degrees from the horizontal. Find the transformation from $\{W\}$ to $\{C\}$.



Example (continued)

- The origin of the camera in the world is ${}^W\mathbf{t}_{cor g} = (0, -5, 3)^T$
- The rotation matrix is ${}^W_C\mathbf{R} = \mathbf{R}_x(-120 \text{ deg})$

$${}^W_C\mathbf{R} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos(-120) & -\sin(-120) \\ 0 & \sin(-120) & \cos(-120) \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & -0.5 & +0.86 \\ 0 & -0.86 & -0.5 \end{bmatrix}$$

- To check to see if this makes sense ... see if the unit axes of the camera are pointing in the right direction in the world

$${}^W_C\mathbf{R} = \left(\begin{pmatrix} {}^W\hat{\mathbf{x}}_C \\ {}^W\hat{\mathbf{y}}_C \\ {}^W\hat{\mathbf{z}}_C \end{pmatrix} \right) \Rightarrow {}^W\hat{\mathbf{x}}_C = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix}, {}^W\hat{\mathbf{y}}_C = \begin{pmatrix} 0 \\ -0.5 \\ -0.86 \end{pmatrix}, {}^W\hat{\mathbf{z}}_C = \begin{pmatrix} 0 \\ 0.86 \\ -0.5 \end{pmatrix}$$

Example (continued)

- The full 4x4 homogeneous transformation matrix from camera to world is

$${}^W_C \mathbf{H} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & -0.5 & 0.86 & -5 \\ 0 & -0.86 & -0.5 & 3 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- But, we actually wanted the transformation from world to camera, so take the inverse

$${}^C_W \mathbf{H} = {}^W_C \mathbf{H}^{-1}$$

1.0000	0	0	0
0	-0.5000	-0.8660	0.0981
0	0.8660	-0.5000	5.8301
0	0	0	1.0000

Python code

```
import math
import numpy as np

def main():
    # Construct transformation from camera to world.
    tc_w = np.array([[0, -5.0, 3.0]]).T
    ax = math.radians(-120) # Convert degrees to radians
    sx = math.sin(ax)
    cx = math.cos(ax)

    Rx = np.array(((1, 0, 0), (0, cx, -sx), (0, sx, cx)))
    R_c_w = Rx # The only rotation is about x

    H_c_w = np.block([[R_c_w, tc_w], [0,0,0,1]]) # Get as 4x4 matrix
    print("H_c_w:"), print(H_c_w)

    # Get transformation from world to camera.
    H_w_c = np.linalg.inv(H_c_w)
    print("H_w_c:"), print(H_w_c)

if __name__ == "__main__":
    main()
```

Summary

- 3D rigid body transformations (i.e., a rotation and translation) can be represented by a single 4x4 homogeneous transformation matrix
- A 3D rotation is represented uniquely by a 3x3 rotation matrix
- 3D rotations can also be represented by XYZ angles (they are easy to understand, but not computationally stable; also the order matters)