

An aerial photograph of the Colorado School of Mines campus. In the foreground, there are several large, modern, light-colored buildings with flat roofs and large windows. A parking lot with a few cars is visible. In the middle ground, there is a large green field, possibly a sports field, surrounded by trees and smaller buildings. In the background, there are rolling hills and mountains under a clear blue sky with a few wispy clouds. The text "Colorado School of Mines" is overlaid in white at the top.

Colorado School of Mines

# Computer Vision

**Professor William Hoff**


Dept of Electrical Engineering & Computer Science

<http://inside.mines.edu/~whoff/>

# Binary Image Processing

# Binary Images

- “Binary” means
  - 0 or 1 values only
  - Or, true/false
- Obtained from
  - Thresholding gray level images
  - Or, the result of feature detectors
- Often want to count or measure shape of 2D binary image regions




# Thresholding


- Convert gray scale image to binary (0s and 1s)
- Bright pixels are mapped to 1s, dark pixels are mapped to 0s
- Need to pick a threshold value



[https://www.bogotobogo.com/Matlab/images/MATLAB\\_DEMO\\_IMAGES/eight.tif](https://www.bogotobogo.com/Matlab/images/MATLAB_DEMO_IMAGES/eight.tif)



Pick, say 200 for  
the threshold



```
B = img > 200 # Creates a boolean image
```

```
from matplotlib import pyplot as plt  
plt.figure() # Matplotlib can display boolean images  
plt.imshow(B, cmap="gray")  
plt.show()
```

# Otsu's Method for Global Thresholding


- Find a threshold that minimizes the variance within groups

$$\sigma_W^2 = P_1 \sigma_1^2 + P_2 \sigma_2^2$$

- where

$$P_1 = \sum_{i=0}^k p_i, \quad P_2 = \sum_{i=k+1}^{L-1} p_i$$

- Intuitively, we want to have each group to be tightly clustered



# OpenCV Thresholding Functions

*# Note - if you have a color image, you need to convert it to grayscale.*

```
gray_img = cv2.cvtColor(bgr_img, cv2.COLOR_BGR2GRAY)
```

*# Do thresholding using a preset threshold. Returns image of type 8-bit unsigned.*

```
_, binary_img = cv2.threshold(gray_img, thresh=127, maxval=255, type=cv2.THRESH_BINARY)
```

Use underscore to ignore the return argument (which is just the threshold we provided)

*# Do thresholding using Otsu's algorithm. The computed threshold value is returned.*

```
thresh, binary_img = cv2.threshold(  
    gray_img, thresh=0, maxval=255, type=cv2.THRESH_BINARY + cv2.THRESH_OTSU)
```

Input threshold is ignored

Note that we can combine flags (see next slide)

# OpenCV Thresholding Flags

Usual type, where lighter values are mapped to white



If you need to map darker values to white



Compute threshold automatically



THRESH_BINARY Python: cv.THRESH_BINARY	$\text{dst}(x, y) = \begin{cases} \text{maxval} & \text{if } \text{src}(x, y) > \text{thresh} \\ 0 & \text{otherwise} \end{cases}$
THRESH_BINARY_INV Python: cv.THRESH_BINARY_INV	$\text{dst}(x, y) = \begin{cases} 0 & \text{if } \text{src}(x, y) > \text{thresh} \\ \text{maxval} & \text{otherwise} \end{cases}$
THRESH_TRUNC Python: cv.THRESH_TRUNC	$\text{dst}(x, y) = \begin{cases} \text{threshold} & \text{if } \text{src}(x, y) > \text{thresh} \\ \text{src}(x, y) & \text{otherwise} \end{cases}$
THRESH_TOZERO Python: cv.THRESH_TOZERO	$\text{dst}(x, y) = \begin{cases} \text{src}(x, y) & \text{if } \text{src}(x, y) > \text{thresh} \\ 0 & \text{otherwise} \end{cases}$
THRESH_TOZERO_INV Python: cv.THRESH_TOZERO_INV	$\text{dst}(x, y) = \begin{cases} 0 & \text{if } \text{src}(x, y) > \text{thresh} \\ \text{src}(x, y) & \text{otherwise} \end{cases}$
THRESH_MASK Python: cv.THRESH_MASK	
THRESH_OTSU Python: cv.THRESH_OTSU	flag, use Otsu algorithm to choose the optimal threshold value
THRESH_TRIANGLE Python: cv.THRESH_TRIANGLE	flag, use Triangle algorithm to choose the optimal threshold value

From documentation for “threshold” at <https://docs.opencv.org>




# Otsu Thresholding Example

- Image is from [https://www.bogotobogo.com/Matlab/images/MATLAB\\_DEMO\\_IMAGES](https://www.bogotobogo.com/Matlab/images/MATLAB_DEMO_IMAGES)



cameraman.tif



Computed threshold = 88.0





# Adaptive Thresholding

- Do a threshold over small windows
- Within each window, compare each pixel to a local mean of the window
- Useful when lighting is uneven

```
binary_img = cv2.adaptiveThreshold(  
    src=gray_img,  
    maxValue=255, # output value where condition met  
    adaptiveMethod=cv2.ADAPTIVE_THRESH_MEAN_C,  
    thresholdType=cv2.THRESH_BINARY, # threshold_type  
    blockSize=51, # neighborhood size (a large odd number)  
    C=-10) # a constant to subtract from mean
```





[https://www.bogotobogo.com/Matlab/images/MATLAB\\_DEMO\\_IMAGES/rice.png](https://www.bogotobogo.com/Matlab/images/MATLAB_DEMO_IMAGES/rice.png)


# Connected Components

- Define adjacency

- 4-adjacent
- 8-adjacent




- Two pixels are connected in S if there is a path between them consisting entirely of pixels in S





- S is a (4- or 8-) connected component (“blob”) if there exists a path between every pair of pixels

- “Labeling” is the process of assigning the same label number to every pixel in a connected component



# Example

- Hand label simple binary image



	1	1				
	1	1				
			1	1		
	1		1	1		
	1		1			1



Binary image


Labeled image (4-connected)


Labeled image (8-connected)

# Example

- Hand label simple binary image



	1	1				
	1	1				
			1	1		
	1		1	1		
	1		1			1

Binary image

	1	1				
	1	1				
			2	2		
	3		2	2		
	3		2			4


Labeled image (4-connected)

	1	1				
	1	1				
			1	1		
	2		1	1		
	2		1			3


Labeled image (8-connected)

# Another example

Binary image



Final labels



- Connected component labeling is fast
- It can be done with only two passes through the image
  - For a nice animation of the two pass algorithm, see <https://iq.opengenus.org/connected-component-labeling/>

# OpenCV Connected Component Labeling

- Find connected components in a binary image

```
# Find connected component labels for all white blobs.
num_white_labels, labels_white_img = cv2.connectedComponents(binary_img)
print("Number of white labels = ", num_white_labels)

# Scale (for display purposes only).
# Min value is set to alpha, max value to beta.
labels_display = cv2.normalize(
    src = labels_white_img, dst = None, alpha = 0, beta = 255,
    norm_type = cv2.NORM_MINMAX, dtype = cv2.CV_8U)


cv2.imshow("Labels image", labels_display)
```

- If you need to find connected components for *black* blobs, complement the binary image


```
# Find connected component labels for all black blobs.
num_black_labels, labels_black_img = cv2.connectedComponents(cv2.bitwise_not(binary_img))
```

# Example


- Labels are displayed as gray levels, 1..N



Input binary image




Note the little white blob!



White labels

Number of white labels = 18



Black labels

Number of black labels = 20



# Binary Image Morphology

- We can “clean up” a binary image, before finding connected components
  - Get rid of tiny regions or holes
- We’ll look at “morphological operations”:
  - Dilation and erosion
  - Opening and closing
- Operations are performed with a “structuring element”  $S$ 
  - $S$  is a small binary image
  - Like a filter mask

# Dilation

- Defined as

$$B \oplus S = \bigcup_{b \in B} S_b$$

- where

- $S_b$  is the structuring element  $S$ , shifted to  $b$

- Procedure

- Sweep  $S$  over  $B$
- Everywhere the origin of  $S$  touches a 1, OR  $S$  with the result image

- Expands regions

1	1	1
1	1	1
1	1	1

$S$

		1		1	
		1	1		

$B$


$B \oplus S$

# Dilation

- Defined as

$$B \oplus S = \bigcup_{b \in B} S_b$$

- where

- $S_b$  is the structuring element  $S$ , shifted to  $b$

- Procedure

- Sweep  $S$  over  $B$
- Everywhere the origin of  $S$  touches a 1, OR  $S$  with the result image

- Expands regions

1	1	1
1	1	1
1	1	1

$S$

		1		1	
		1	1		

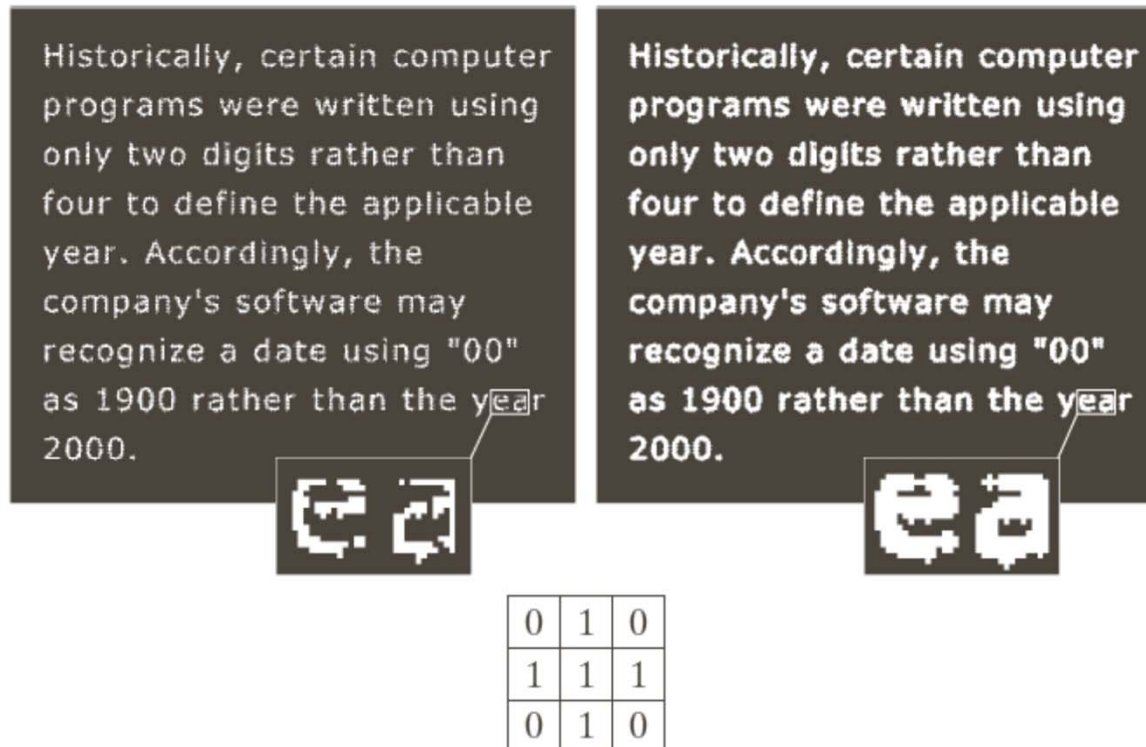

$B$

	1	1	1	1	1
	1	1	1	1	1
	1	1	1	1	1
	1	1	1	1	

$B \oplus S$

# Dilation

- Dilation fills in holes, thickens thin parts, grows object



a b c

**FIGURE 9.7**  
 (a) Sample text of poor resolution with broken characters (see magnified view).  
 (b) Structuring element.  
 (c) Dilation of (a) by (b). Broken segments were joined.

# Erosion

- Defined as

$$B \ominus S = \{b \mid b + s \in B, \forall s \in S\}$$

- Procedure

- Sweep S over B
- Everywhere S is completely contained in B, output a 1 at the origin of S

- Shrinks regions

1	1	1
1	1	1
1	1	1

S

	1	1	1	1	1
	1	1	1	1	1
	1	1	1	1	1
	1	1	1	1	

B


$B \ominus S$

# Erosion

- Defined as

$$B \ominus S = \{b \mid b + s \in B, \forall s \in S\}$$

- Procedure

- Sweep S over B
- Everywhere S is completely contained in B, output a 1 at the origin of S

- Shrinks regions

1	1	1
1	1	1
1	1	1

S

	1	1	1	1	1
	1	1	1	1	1
	1	1	1	1	1
	1	1	1	1	

B

		1	1	1	
		1	1		

$B \ominus S$

# Openings and Closings

- Opening
  - Erosion followed by dilation
  - Eliminate small regions and projections

$$B \circ S = (B \ominus S) \oplus S$$

- Closing
  - Dilation followed by erosion
  - Fill in small holes and gaps



$$B \bullet S = (B \oplus S) \ominus S$$

*Advantage of openings and closings: doesn't change size of large regions*





# Example - opening



*Eliminate small  
regions and  
projections*



erosion




dilation




$$A \circ B = (A \ominus B) \oplus B$$

# Example - closing



*Fill in small  
holes and gaps*



$A$





dilation



$A \oplus B$

erosion



$$A \cdot B = (A \oplus B) \ominus B$$

# OpenCV

- To create a structuring element, or “kernel”

*# Create a 5x5 square box filter.*

```
kernel = np.ones((5, 5), np.uint8)
```

```
print(kernel)
```

*# Or, make a disk, if we want something more well rounded.*

```
kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (5, 5))
```

- Opening

```
filtered_img = cv2.morphologyEx(binary_img, cv2.MORPH_OPEN, kernel)
```

- Closing

```
filtered_img = cv2.morphologyEx(binary_img, cv2.MORPH_CLOSE, kernel)
```

# Example – real image

- Segment all the dark regions in the lower half of the image “bag.png”
  - Namely, generate a binary (or “logical”) image which is white (1) in the regions of interest, and black (0) elsewhere
  - Want:
    - No gaps in the regions
    - No extraneous white pixels in the background
- Then do connected component labeling on these regions




Image “bag.png” from  
[https://www.bogotobogo.com/  
Matlab/images/MATLAB\\_DEMO\\_IMAGES/](https://www.bogotobogo.com/Matlab/images/MATLAB_DEMO_IMAGES/)

*For this example,  
ignore the upper  
half of the image*

# Approach

```
# Complement the binary image so that the black regions are now white.  
binary_img = cv2.bitwise_not(binary_img)
```

```
# Clean up using opening + closing.
```

```
ksize = 5
```

```
kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (ksize, ksize))
```

```
binary_img = cv2.morphologyEx(binary_img, cv2.MORPH_OPEN, kernel)
```

```
binary_img = cv2.morphologyEx(binary_img, cv2.MORPH_CLOSE, kernel)
```

```
cv2.imshow("Filtered image", binary_img)
```


```
# Find connected component labels for all white blobs.
```

```
num_white_labels, labels_white_img = cv2.connectedComponents(binary_img)
```


*You can find the value of  
"ksize" experimentally*

*You can use a different  
structuring element for opening  
and closing (not typical)*


Number of white labels = 27




Original "bag.png" image



Binary image after Otsu thresholding



After morphological operations



After connected component labeling. Each gray level indicates a different label of a foreground object

# Region Properties

- Area


$$A = \sum_{(r,c) \in R} 1$$

- Centroid

$$\bar{r} = \frac{1}{A} \sum_{(r,c) \in R} r, \quad \bar{c} = \frac{1}{A} \sum_{(r,c) \in R} c$$

- Bounding box

- The smallest rectangle containing the region
- Can be specified by
  - The location of the upper left corner
  - The width and height



# OpenCV - computing region properties

- Get connected components and region properties


```
num_labels, labels_img, stats, centroids = cv2.connectedComponentsWithStats(binary_img)
```

- Print centroids and areas

```
for n in range(num_labels):  
    xc, yc = centroids[n]  
    area = stats[n, cv2.CC_STAT_AREA]  
    print("Region %d, centroid: (%f,%f), area = %d" % (n, xc, yc, area))
```

- Display bounding boxes


```
bgr_image_display = cv2.cvtColor(binary_img, cv2.COLOR_GRAY2BGR)  
for stat, centroid in zip(stats, centroids):  
    x0 = stat[cv2.CC_STAT_LEFT]  
    y0 = stat[cv2.CC_STAT_TOP]  
    w = stat[cv2.CC_STAT_WIDTH]  
    h = stat[cv2.CC_STAT_HEIGHT]  
    bgr_image_display = cv2.rectangle(  
        img=bgr_image_display, pt1=(x0, y0), pt2=(x0 + w, y0 + h),  
        color=(0, 0, 255), thickness=1)  
cv2.imshow("boxes", bgr_image_display)
```







# Concentric contrasting circle (CCC) target

- The target is a white ring surrounding a black dot
- This feature is fairly unique in the image, because the centroid of the white ring will coincide with the centroid of the black dot
- You can automatically find the target by finding a white region whose centroid coincides with the centroid of a black region



# CCC targets (continued)

- For more discrimination power, you can also place constraints on the binary regions, e.g.,
  - The white area must be greater than the black area)
  - The white bounding box must enclose the black bounding box



- Also, sometimes adaptive thresholding is better than global thresholding; especially when the lighting varies across the image




# Overall strategy for finding CCC targets...

- Threshold the image
- Clean up the image using opening and closing
- Find connected components
- Find white blobs
- Complement the image and repeat the process to find black blobs
- Check every possible pair of white and black blobs...

```
for each white blob
  for each black blob
    Get centroid of white blob
    Get centroid of black blob
    if distance between centroids < thresh
      (Potentially also check if one bounding box encloses the other)
      We have a possible CCC! draw a crosshair at its centroid,
      or draw a rectangle around its bounding box
    end if
  end for
end for
```

# Example

*Here, I used a threshold of 3.0 pixels to check if centroids coincide*



# Fun fact – circle targets in space!

